

Typeclasses

Garbage Collection

Feb 2021

Takashi Idobe

-
- Imagine you have your own room, and you have to take out the trash every once in a while.
 - That's basically garbage collection.
-

Strategies

- No GC
 - Reference Counting
 - Mark-Sweep
 - Multi-Space
-

No GC

- We allocate a chunk of memory up front, and when we use it all up, we crash.
 - At the end, we throw everything away.
-

Visualization

- Black = Memory not in use
 - Green + Yellow = Memory in Use
 - Gray = Memory not in use for long time.
-
-

Usage

- High Frequency Traders do this a lot.

Pros

- Fast and Simple

Cons

- We can't free anything.
 - Memory Fragmentation.
-

Reference Counting

- Keep track of how many times you're using a resource.
 - When it drops to 0, free it.
-
-

Usage

- Obj-C, Swift, Rust, C++

Pros

- Easy to retrofit
- Simple?

Cons

- Can't handle cyclic structures (like Graphs)
 - Counters aren't thread-safe, so perf hit.
 - Can trigger large GC pauses on frequently referenced data.
-

Mark Sweep Collector

- Every once in a while, pause the program's execution, and see if objects are reachable.
 - If reachable, mark as reachable.
 - If not reachable, reclaim.
-
-

Usage

- Lisp, Java, etc.

Pros

- Less overhead than RC.
- No counting required

Cons

- Fragmentation
 - Can trigger large GC pauses when there's a large amount of data.
-

Mark Compact

- Mark-Sweep, but compacts memory after sweep.
 - This requires an extra sweep to figure out where to put items.
-
-

Usage

- Lisp, Java, etc.

Pros

- Less Fragmentation
- Better Performance?

Cons

- GC Pauses are longer than Mark-Sweep.
-

Copying Collector

- People figured out that young objects tend to be freed more than older objects.
 - Create an area for young objects and one for old objects.
 - Check the young area more often, and move young objects to the old area when some time has passed.
-
-

Usage

- Lisp, Java, etc.

Pros

- Fewer Sweeps
- Less Fragmentation

Cons

- Memory might take more time to be freed.
- Requires moving memory around.

GC Throughput




Figure 1: GC Speed

Pause Duration

Takeaways

- GC performance can vary wildly based on what GC you choose.
- Pick one that suits your workload.
- It's hard to make a general purpose GC.




Figure 2: GC Pause